

Algorithms in MapReduce

Krzysztof Dembczyński

Intelligent Decision Support Systems Laboratory (IDSS)
Poznań University of Technology, Poland



Intelligent Decision Support Systems
Master studies, second semester
Academic year 2015/16 (summer course)

Review of the previous lectures

- Mining of massive datasets.
- Evolution of database systems.
- Dimensional modeling.
- ETL and OLAP systems.
- Multidimensional queries.
- Processing of very large data.
- Nearest neighbor search.
- Data partitioning and MapReduce:
 - ▶ Data partitioning.
 - ▶ The overall idea of the MapReduce paradigm.
 - ▶ Hadoop implementation.

Outline

- 1 Motivation
- 2 Algorithms in Map-Reduce
- 3 Extensions to MapReduce
- 4 Summary

Outline

- 1 Motivation
- 2 Algorithms in Map-Reduce
- 3 Extensions to MapReduce
- 4 Summary

Algorithms in Map-Reduce

- How to implement fundamental algorithms in MapReduce?
 - ▶ Matrix-vector multiplication.
 - ▶ Relational-Algebra Operations.
 - ▶ Matrix multiplication.

Outline

- 1 Motivation
- 2 Algorithms in Map-Reduce**
- 3 Extensions to MapReduce
- 4 Summary

Matrix-vector Multiplication

- Let A to be large $n \times m$ matrix, and x a long vector of size m .
- The matrix-vector multiplication is defined as:

Matrix-vector Multiplication

- Let \mathbf{A} to be large $n \times m$ matrix, and \mathbf{x} a long vector of size m .
- The matrix-vector multiplication is defined as:

$$\mathbf{Ax} = \mathbf{v},$$

where $\mathbf{v} = (v_1, \dots, v_n)$ and

$$v_i = \sum_{j=1}^m a_{ij}x_j.$$

Matrix-vector multiplication

- Let us first assume that m is large, but not so large that vector \mathbf{x} cannot fit in main memory, and be part of the input to every Map task.
- The matrix \mathbf{A} is stored with explicit coordinates, as a triple (i, j, a_{ij}) .
- We also assume the position of element x_j in the vector \mathbf{x} will be stored in the analogous way.

Matrix-vector multiplication

- **Map:**

Matrix-vector multiplication

- **Map:** each map task will take the entire vector x and a chunk of the matrix A . From each matrix element a_{ij} it produces the key-value pair $(i, a_{ij}x_j)$. Thus, all terms of the sum that make up the component v_i of the matrix-vector product will get the same key.

Matrix-vector multiplication

- **Map:** each map task will take the entire vector x and a chunk of the matrix A . From each matrix element a_{ij} it produces the key-value pair $(i, a_{ij}x_j)$. Thus, all terms of the sum that make up the component v_i of the matrix-vector product will get the same key.
- **Reduce:**

Matrix-vector multiplication


- **Map:** each map task will take the entire vector \mathbf{x} and a chunk of the matrix \mathbf{A} . From each matrix element a_{ij} it produces the key-value pair $(i, a_{ij}x_j)$. Thus, all terms of the sum that make up the component v_i of the matrix-vector product will get the same key.
- **Reduce:** a reduce task has simply to sum all the values associated with a given key i . The result will be a pair (i, v_i) where:

$$v_i = \sum_{j=1}^m a_{ij}x_j.$$

Matrix-Vector Multiplication with Large Vector v

Matrix-Vector Multiplication with Large Vector v

- Divide the matrix into vertical stripes of equal width and divide the vector into an equal number of horizontal stripes, of the same height.



- The i th stripe of the matrix multiplies only components from the i th stripe of the vector.
- Thus, we can divide the matrix into one file for each stripe, and do the same for the vector.

Matrix-Vector Multiplication with Large Vector v

- Each Map task is assigned a chunk from one the stripes of the matrix and gets the entire corresponding stripe of the vector.
- The Map and Reduce tasks can then act exactly as in the case where Map tasks get the entire vector.

Relational-Algebra Operations

Example (Relation **Links**)

From	To
url1	url2
url1	url3
url2	url3
url2	url4
...	...

Relational-Algebra Operations

- Selection
- Projection
- Union, Intersection, and Difference
- Natural Join
- Grouping and Aggregation

Relational-Algebra Operations

- R, S - relation
- t, t' - a tuple
- \mathcal{C} - a condition of selection
- A, B, C - subset of attributes
- a, b, c - attribute values for a given subset of attributes

Selection

- **Map:**

Selection

- **Map:** For each tuple t in R , test if it satisfies \mathcal{C} . If so, produce the key-value pair (t, t) . That is, both the key and value are t .
- **Reduce:**

Selection

- **Map:** For each tuple t in R , test if it satisfies \mathcal{C} . If so, produce the key-value pair (t, t) . That is, both the key and value are t .
- **Reduce:** The Reduce function is the identity. It simply passes each key-value pair to the output.

Projection

Projection

- **Map:**

Projection

- **Map:** For each tuple t in R , construct a tuple t' by eliminating from t those components whose attributes are not in A . Output the key-value pair (t', t') .
- **Reduce:**

Projection

- **Map:** For each tuple t in R , construct a tuple t' by eliminating from t those components whose attributes are not in A . Output the key-value pair (t', t') .
- **Reduce:** For each key t' produced by any of the Map tasks, there will be one or more key-value pairs (t', t') . The Reduce function turns $(t', [t', t', \dots, t'])$ into (t', t') , so it produces exactly one pair (t', t') for this key t' .

Union

- **Map:**

Union

- **Map:** Turn each input tuple t either from relation R or S into a key-value pair (t, t) .
- **Reduce:**

Union

- **Map:** Turn each input tuple t either from relation R or S into a key-value pair (t, t) .
- **Reduce:** Associated with each key t there will be either one or two values. Produce output (t, t) in either case.

Intersection

- **Map:**

Intersection

- **Map:** Turn each input tuple t either from relation R or S into a key-value pair (t, t) .
- **Reduce:**

Intersection

- **Map:** Turn each input tuple t either from relation R or S into a key-value pair (t, t) .
- **Reduce:** If key t has value list $[t, t]$, then produce (t, t) . Otherwise, produce nothing.

Minus

- **Map:**

Minus

- **Map:** For a tuple t in R , produce key-value pair $(t, \text{name}(R))$, and for a tuple t in S , produce key-value pair $(t, \text{name}(S))$.
- **Reduce:**

Minus

- **Map:** For a tuple t in R , produce key-value pair $(t, \text{name}(R))$, and for a tuple t in S , produce key-value pair $(t, \text{name}(S))$.
- **Reduce:** For each key t , do the following.
 - 1 If the associated value list is $[\text{name}(R)]$, then produce (t, t) .
 - 2 If the associated value list is anything else, which could only be $[\text{name}(R), \text{name}(S)]$, $[\text{name}(S), \text{name}(R)]$, or $[\text{name}(S)]$, produce nothing.

Natural Join

- Let us assume that we join relation $R(A, B)$ with relation $S(B, C)$ that share the same attribute B .
- **Map:**

Natural Join

- Let us assume that we join relation $R(A, B)$ with relation $S(B, C)$ that share the same attribute B .
- **Map:** For each tuple (a, b) of R , produce the key-value pair $(b, (\text{name}(R), a))$. For each tuple (b, c) of S , produce the key-value pair $(b, (\text{name}(S), c))$.
- **Reduce:**

Natural Join

- Let us assume that we join relation $R(A, B)$ with relation $S(B, C)$ that share the same attribute B .
- **Map:** For each tuple (a, b) of R , produce the key-value pair $(b, (\text{name}(R), a))$. For each tuple (b, c) of S , produce the key-value pair $(b, (\text{name}(S), c))$.
- **Reduce:** Each key value b will be associated with a list of pairs that are either of the form $(\text{name}(R), a)$ or $(\text{name}(S), c)$. Construct all pairs consisting of one with first component $\text{name}(R)$ and the other with first component S , say $(\text{name}(R), a)$ and $(\text{name}(S), c)$. The output for key b is $(b, [(a1, b, c1), (a2, b, c2), \dots])$, that is, b associated with the list of tuples that can be formed from an R -tuple and an S -tuple with a common b value.

Grouping and Aggregation

- Let assume that we group a relation $R(A, B, C)$ by attributes A and aggregate values of B .
- **Map:**

Grouping and Aggregation

- Let assume that we group a relation $R(A, B, C)$ by attributes A and aggregate values of B .
- **Map**: For each tuple (a, b, c) produce the key-value pair (a, b) .
- **Reduce**:

Grouping and Aggregation

- Let assume that we group a relation $R(A, B, C)$ by attributes A and aggregate values of B .
- **Map**: For each tuple (a, b, c) produce the key-value pair (a, b) .
- **Reduce**: Each key a represents a group. Apply the aggregation operator θ to the list $[b_1, b_2, \dots, b_n]$ of B -values associated with key a . The output is the pair (a, x) , where x is the result of applying θ to the list. For example, if θ is SUM, then $x = b_1 + b_2 + \dots + b_n$, and if θ is MAX, then x is the largest of b_1, b_2, \dots, b_n .

Matrix Multiplication

- If M is a matrix with element m_{ij} in row i and column j , and N is a matrix with element n_{jk} in row j and column k , then the product:

$$P = MN$$

is the matrix P with element p_{ik} in row i and column k , where:

$$p_{ik} =$$

Matrix Multiplication

- If M is a matrix with element m_{ij} in row i and column j , and N is a matrix with element n_{jk} in row j and column k , then the product:

$$P = MN$$

is the matrix P with element p_{ik} in row i and column k , where:

$$p_{ik} = \sum_j m_{ij} n_{jk}$$

Matrix Multiplication

- We can think of a matrix M and N as a relation with three attributes: the row number, the column number, and the value in that row and column, i.e.,:

$$M(I, J, V) \quad \text{and} \quad N(J, K, W)$$

with the following tuples, respectively:

$$(i, j, m_{ij}) \quad \text{and} \quad (j, k, n_{jk}).$$

- In case of sparsity of M and N , this relational representation is very efficient in terms of space.
- The product MN is almost a natural join followed by grouping and aggregation.

Matrix Multiplication

Matrix Multiplication

- **Map:**

Matrix Multiplication

- **Map:** Send each matrix element m_{ij} to the key value pair:

$$(j, (M, i, m_{ij})) .$$

Analogously, send each matrix element n_{jk} to the key value pair:

$$(j, (N, k, n_{jk})) .$$

- **Reduce:**

Matrix Multiplication

- **Map:** Send each matrix element m_{ij} to the key value pair:

$$(j, (M, i, m_{ij})) .$$

Analogously, send each matrix element n_{jk} to the key value pair:

$$(j, (N, k, n_{jk})) .$$

- **Reduce:** For each key j , examine its list of associated values. For each value that comes from M , say (M, i, m_{ij}) , and each value that comes from N , say (N, k, n_{jk}) , produce the tuple

$$(i, k, v = m_{ij}n_{jk}),$$

The output of the Reduce function is a key j paired with the list of all the tuples of this form that we get from j :

$$(j, [(i_1, k_1, v_1), (i_2, k_2, v_2), \dots, (i_p, k_p, v_p)]) .$$

Matrix Multiplication

Matrix Multiplication

- **Map:**

Matrix Multiplication

- **Map:** From the pairs that are output from the previous Reduce function produce p key-value pairs:

$$((i_1, k_1), v_1), ((i_2, k_2), v_2), \dots, ((i_p, k_p), v_p) .$$

- **Reduce:**

Matrix Multiplication

- **Map:** From the pairs that are output from the previous Reduce function produce p key-value pairs:

$$((i_1, k_1), v_1), ((i_2, k_2), v_2), \dots, ((i_p, k_p), v_p).$$

- **Reduce:** For each key (i, k) , produce the sum of the list of values associated with this key. The result is a pair

$$((i, k), v),$$

where v is the value of the element in row i and column k of the matrix

$$P = MN.$$

Matrix Multiplication with One Map-Reduce Step

- **Map:**

Matrix Multiplication with One Map-Reduce Step

- **Map:** For each element m_{ij} of M , produce a key-value pair

$$((i, k), (M, j, m_{ij})),$$

for $k = 1, 2, \dots$, up to the number of columns of N .

Also, for each element n_{jk} of N , produce a key-value pair

$$((i, k), (N, j, n_{jk})),$$

for $i = 1, 2, \dots$, up to the number of rows of M .

Matrix Multiplication with One Map-Reduce Step

- **Reduce:**

Matrix Multiplication with One Map-Reduce Step

- **Reduce:** Each key (i, k) will have an associated list with all the values

$$(M, j, m_{ij}) \quad \text{and} \quad (N, j, n_{jk}),$$

for all possible values of j . We connect the two values on the list that have the same value of j , for each j :

- ▶ We sort by j the values that begin with M and sort by j the values that begin with N , in separate lists,
- ▶ The j th values on each list must have their third components, m_{ij} and n_{jk} extracted and multiplied,
- ▶ Then, these products are summed and the result is paired with (i, k) in the output of the Reduce function.

Outline

- 1 Motivation
- 2 Algorithms in Map-Reduce
- 3 Extensions to MapReduce**
- 4 Summary

Extensions to MapReduce

Extensions to MapReduce

- Workflow systems, like Clustera or Hyracks, extend map-reduce from the simple two-step workflow (the Map function feeds the Reduce function) to any collection of functions, with an acyclic graph representing workflow among the functions.

Extensions to MapReduce

- Workflow systems, like Clustera or Hyracks, extend map-reduce from the simple two-step workflow (the Map function feeds the Reduce function) to any collection of functions, with an acyclic graph representing workflow among the functions.
- Pregel is a distributed programming framework for graph algorithms
 - ▶ Pregel views its data as a graph.
 - ▶ Each node of the graph corresponds roughly to a task
 - ▶ Each graph node generates output messages that are destined for other nodes of the graph, and each graph node processes the inputs it receives from other nodes.

Extensions to MapReduce

- Workflow systems, like Clustera or Hyracks, extend map-reduce from the simple two-step workflow (the Map function feeds the Reduce function) to any collection of functions, with an acyclic graph representing workflow among the functions.
- Pregel is a distributed programming framework for graph algorithms
 - ▶ Pregel views its data as a graph.
 - ▶ Each node of the graph corresponds roughly to a task
 - ▶ Each graph node generates output messages that are destined for other nodes of the graph, and each graph node processes the inputs it receives from other nodes.
- PIG and Hive: the former is an implementation of relational algebra on top of Hadoop, while the latter implements a restricted form of SQL on top of Hadoop.
- And many others ...

Outline

- 1 Motivation
- 2 Algorithms in Map-Reduce
- 3 Extensions to MapReduce
- 4 Summary

Summary

- Algorithms in MapReduce:
 - ▶ Matrix-vector multiplication.
 - ▶ Relational-Algebra Operations.
 - ▶ Matrix multiplication.
- Extensions of MapReduce

Bibliography

- A. Rajaraman and J. D. Ullman. *Mining of Massive Datasets*.
Cambridge University Press, 2011
<http://infolab.stanford.edu/~ullman/mmds.html>
- J.Lin and Ch. Dyer. *Data-Intensive Text Processing with MapReduce*.
Morgan and Claypool Publishers, 2010
<http://lintool.github.com/MapReduceAlgorithms/>
- Ch. Lam. *Hadoop in Action*.
Manning Publications Co., 2011